

# sample

December 2, 2025

## 0.1 Part 1: Dataset Simulation

[1]: # Library imports

```
import pandas as pd
import numpy as np
from datetime import datetime
import random
from google import genai
from google.genai import types
from collections import Counter
import os
import math
import time
import json

import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, LabelEncoder

from dotenv import load_dotenv
load_dotenv()

output_filename = 'data/sim_customers_seed_42.csv'
```

[6]: # Helper: robustly convert various SDK response types to plain string for JSON parsing

```
def _extract_text(obj):
    """Recursively extract text from SDK response objects.
    Handles str/bytes, lists, objects with .text/.content/.output, pydantic
    models, response.candidates, and fallbacks to str().
    """
    if obj is None:
        return ""
    if isinstance(obj, str):
```

```

    return obj
if isinstance(obj, (bytes, bytearray)):
    try:
        return obj.decode('utf-8')
    except Exception:
        return str(obj)
# If object has 'parts' (GenAI SDK), extract part contents
try:
    if hasattr(obj, 'parts') and obj.parts:
        # join all part contents
        return "\n".join(_extract_text(p) for p in obj.parts)
except Exception:
    pass
# If it's a list or tuple, join items
if isinstance(obj, (list, tuple)):
    return "\n".join(_extract_text(x) for x in obj)
# Common nested attributes
for attr in ("text", "content", "output", "value", "data"):
    try:
        val = getattr(obj, attr, None)
    except Exception:
        val = None
    if val:
        return _extract_text(val)
# If object is a dict-like
try:
    if isinstance(obj, dict):
        return json.dumps(obj)
except Exception:
    pass
# Try pydantic/BaseModel .json() or .dict()
try:
    if hasattr(obj, 'json') and callable(obj.json):
        return obj.json()
except Exception:
    pass
try:
    if hasattr(obj, 'dict') and callable(obj.dict):
        return json.dumps(obj.dict())
except Exception:
    pass
# Fallback to string
return str(obj)

```

[26]: # Methods used in the simulation

```
# Set random seed for reproducibility
```

```

np.random.seed(42)
random.seed(42)

TRANSACTION_CATEGORIES = ["groceries", "utilities", "entertainment", "dining", ↴
    "transport", "healthcare", "shopping", "education", "rent", "salary", ↴
    "transfer"]

def generate_feedback_gemini_batch(summaries, n_requests=10, model="gemini-2.5-flash", max_retries=3, retry_delay=5, batch_delay_seconds=4):
    """Send summaries to Gemini in n_requests batches and return feedback map.

    This function WILL raise if the GenAI SDK is missing, the GOOGLE_API_KEY env var
    is not set, or if Gemini fails to produce valid JSON for any batch after retries.

    summaries: list of dicts, each must include 'customer_id' and other summary fields.
    n_requests: desired number of Gemini calls (e.g., 10 to respect quotas).
    batch_delay_seconds: seconds to sleep between successful batch calls to avoid bursts.

    Returns: {customer_id: feedback_text}
    """

    # Strict preconditions for Gemini-only mode
    if 'genai' not in globals() or genai is None or 'types' not in globals() or types is None:
        raise RuntimeError("Google GenAI SDK not available. Install the official SDK before enabling Gemini.")

    GOOGLE_API_KEY = os.getenv("GOOGLE_API_KEY")
    if not GOOGLE_API_KEY:
        raise ValueError("GOOGLE_API_KEY not found in environment variables; required for Gemini mode.")

    client = genai.Client(api_key=GOOGLE_API_KEY)

    feedback_map = {}
    total = len(summaries)
    if total == 0:
        return feedback_map

    batch_size = max(1, math.ceil(total / n_requests))

    for i in range(0, total, batch_size):
        batch = summaries[i:i+batch_size]
        # Build strict JSON-response prompt

```

```

        items_text = []
        for s in batch:
            items_text.append(f"ID: {s['customer_id']} | credits: {s['total_credits']:.2f} | debits: {s['total_debits']:.2f} | end_balance: {s['ending_balance']:.2f} | top: {s['top_category']}")
        prompt_body = "\n".join(items_text)
        prompt = (
            "You are given multiple short customer monthly summaries (one per line in the format 'ID: <id> | ...').\n"
            "For each summary, produce a JSON array of objects with fields {\"customer_id\": <id>, \"feedback\": <short advice sentence>} and NOTHING ELSE.\n"
            "Keep each feedback within 10 words and make them varied and realistic. Example: \"I plan to increase my savings next month.\n"
            "Respond with valid JSON only (no extra commentary).\n\n"
            f"Summaries:\n{prompt_body}"
        )

    attempt = 0
    resp_text = None
    last_exception = None
    while attempt < max_retries:
        try:
            response = client.models.generate_content(
                model=model,
                contents=[prompt],
                config=types.GenerateContentConfig(
                    system_instruction="You are generating short customer feedback notes for a synthetic banking dataset.",
                    thinking_config=types.ThinkingConfig(thinking_budget=-1)
                )
            )
            # Extract returned text conservatively
            resp_text = _extract_text(response)
            if resp_text:
                break
            else:
                raise RuntimeError('Empty response from Gemini')
        except Exception as e:
            last_exception = e
            attempt += 1
            wait = retry_delay * (2 ** (attempt - 1))
            # If the exception contains retry info, prefer that
            # (best-effort)
            try:

```

```

# attempt to read retryDelay from exception message or
object
        # Not all exceptions expose this; this is best-effort
        print(f"Gemini batch request failed (attempt {attempt}/
{max_retries}): {e}. Retrying in {wait}s...")
    except Exception:
        pass
    time.sleep(wait)

if resp_text is None:
    # unrecoverable for this batch
    raise RuntimeError(f"Gemini failed for batch starting at index {i}_
after {max_retries} attempts." from last_exception

# Parse JSON exactly; be strict in Gemini-only mode
# Clean markdown code blocks
if resp_text.startswith('```json') and resp_text.endswith('```'):
    resp_text = resp_text[7:-3].strip()
elif resp_text.startswith('`') and resp_text.endswith('`'):
    resp_text = resp_text[3:-3].strip()
parsed = None
try:
    parsed = json.loads(resp_text)
except Exception:
    # attempt to extract JSON substring
    start = resp_text.find('[')
    end = resp_text.rfind(']')
    if start != -1 and end != -1 and end > start:
        parsed = json.loads(resp_text[start:end+1])

if not isinstance(parsed, list):
    raise ValueError(f"Gemini response for batch starting at index {i}_
did not parse to a JSON list.\nResponse:\n{resp_text}")

for obj in parsed:
    cid = obj.get('customer_id') or obj.get('ID') or obj.get('id')
    fb = obj.get('feedback') or obj.get('note') or obj.get('text')
    if not cid or not fb:
        raise ValueError(f"Invalid object in Gemini JSON output: {obj}")
    feedback_map[str(cid)] = str(fb)

# Sleep a short while between batches to avoid quota bursts
time.sleep(batch_delay_seconds)

return feedback_map

```

```

# Replace daily series to *require* Gemini use and map feedbacks strictly
def generate_customer_dataset(num_customers=500, seed=42, gemini_requests=10, ✉
    ↪batch_delay_seconds=4):
    np.random.seed(seed)
    random.seed(seed)

    rows = []
    month_end_date = datetime(2025, 1, 31).date()

    customer_summaries = []

    for cid in range(1, num_customers + 1):
        customer_id = f"CUST{str(cid).zfill(5)}"
        age = int(np.clip(np.random.normal(40, 12), 18, 80))
        monthly_income = float(np.round(np.clip(np.random.normal(4000, 2000), ✉
            ↪500, None), 2))
        occupation = random.choice(["Engineer", "Teacher", "Healthcare", ✉
            ↪"Retail", "Finance", "Student", "Self-employed", "Unemployed", "Manager", ✉
            ↪"Technician"])
        starting_balance = float(np.round(monthly_income * np.random.uniform(0. ✉
            ↪2, 4.0) + np.random.uniform(-500, 500), 2))
        credit_score = int(np.clip(np.random.normal(680, 50), 300, 850))
        invest_pct = float(np.round(np.random.normal(0.8, 1.2), 2))

        # Simulate monthly transactions
        balance = starting_balance
        monthly_credits_total = 0.0
        monthly_debits_total = 0.0
        categories_counter = Counter()

        # Salary credit
        salary = monthly_income
        monthly_credits_total += salary
        balance += salary
        categories_counter.update(["salary"])

        # Random transfers
        n_transfers = np.random.poisson(2)
        for _ in range(n_transfers):
            transfer = float(np.random.exponential(scale=monthly_income / 10.0))
            monthly_credits_total += transfer
            balance += transfer
            categories_counter.update(["transfer"])

        # Investment gains
        monthly_invest_gain = (balance * max(0.0, invest_pct) / 100.0)
        monthly_credits_total += monthly_invest_gain

```

```

balance += monthly_invest_gain

# Debits
n_txns = np.random.poisson(30) # approx daily
for _ in range(n_txns):
    cat = random.choice(TRANSACTION_CATEGORIES[:-3])
    amt = float(np.random.exponential(scale=monthly_income / 120.0)) +_
        np.random.uniform(5, 50)
    monthly_debits_total += amt
    balance -= amt
    categories_counter.update([cat])

# Big spends
n_big = np.random.poisson(3)
for _ in range(n_big):
    big_cat = random.choice(["shopping", "healthcare", "education",_
        "rent"])
    big_amt = float(np.random.exponential(scale=monthly_income / 4.0))
    monthly_debits_total += big_amt
    balance -= big_amt
    categories_counter.update([big_cat])

ending_balance = float(np.round(balance, 2))

# Credit score change
credit_score_change = np.random.choice([-10, -5, 5, 10])
final_credit_score = int(np.clip(credit_score + credit_score_change,_
    -300, 850))

top_category = categories_counter.most_common(1)[0][0] if_
categories_counter else None

row = {
    "customer_id": customer_id,
    "date": month_end_date.isoformat(),
    "age": age,
    "occupation": occupation,
    "monthly_income": monthly_income,
    "monthly_credit": float(round(monthly_credits_total, 2)),
    "monthly_debit": float(round(monthly_debits_total, 2)),
    "starting_balance": starting_balance,
    "ending_balance": ending_balance,
    "monthly_invest_gain": float(round(monthly_invest_gain, 2)),
    "monthly_invest_return_pct": invest_pct,
    "initial_credit_score": credit_score,
    "final_credit_score": final_credit_score,
    "transaction_category_major": top_category,
}

```

```

        "customer_feedback": None
    }
    rows.append(row)

    summary = {
        "customer_id": customer_id,
        "total_credits": monthly_credits_total,
        "total_debits": monthly_debits_total,
        "ending_balance": ending_balance,
        "top_category": top_category
    }
    customer_summaries.append(summary)

# Generate feedbacks using Gemini batches (Gemini-only mode)
feedback_map = generate_feedback_gemini_batch(customer_summaries, n_requests=gemini_requests, batch_delay_seconds=batch_delay_seconds)

# Ensure every customer got a Gemini feedback
missing = [s['customer_id'] for s in customer_summaries if s['customer_id'] not in feedback_map]
if missing:
    raise RuntimeError(f"Gemini did not return feedback for the following customers: {missing[:10]}{'...' if len(missing)>10 else ''}")

# Attach feedbacks
for r in rows:
    cid = r['customer_id']
    fb = feedback_map[cid]
    r['customer_feedback'] = fb

df = pd.DataFrame(rows)

# Inject missingness and noise
inc_idx = df.sample(frac=0.02, random_state=seed).index
df.loc[inc_idx, "monthly_income"] = np.nan

txn_idx = df.sample(frac=0.03, random_state=seed+1).index
df.loc[txn_idx, ["monthly_credit", "monthly_debit"]] = np.nan

cs_idx = df.sample(frac=0.01, random_state=seed+2).index
df.loc[cs_idx, "final_credit_score"] = df.loc[cs_idx, "final_credit_score"].apply(lambda x: int(x + np.random.choice([-500, 400])))

fb_idx = df[df["customer_feedback"].notnull()].sample(frac=0.05, random_state=seed+3).index
df.loc[fb_idx, "customer_feedback"] = None

```

```

ol_idx = df.sample(n=12, random_state=seed+4).index
df.loc[ol_idx, "monthly_debit"] = df.loc[ol_idx, "monthly_debit"] * np.
    ↪random.uniform(5, 25, size=len(ol_idx))

col_order = [
    "customer_id", "date", "age", "occupation", "monthly_income",
    "monthly_credit", "monthly_debit", "starting_balance", "ending_balance",
    "monthly_invest_gain", "monthly_invest_return_pct", ↪
    ↪"initial_credit_score",
    "final_credit_score", "transaction_category_major", "customer_feedback"
]
df = df[col_order]

return df

def save_simulation(df: pd.DataFrame, path: str):
    os.makedirs(os.path.dirname(path), exist_ok=True)
    df.to_csv(path, index=False)

```

[ ]: # Configuration and run

```

NUM_CUSTOMERS = 500
SEED = 42

print("Generating synthetic banking customer dataset...")
df = generate_customer_dataset(num_customers=NUM_CUSTOMERS, seed=SEED)

print(f"\n Dataset generated successfully! Rows: {len(df)} Customers: ↪
    ↪{df['customer_id'].nunique()}")
print(f"Date: {df['date'].iloc[0]}")

# Show missing data summary
print("\nMissing data summary:")
print(df.isnull().sum())

# Show sample of feedbacks
print("\nSample feedbacks (first 10 customers):")
sample_feedbacks = df[['customer_id', 'customer_feedback']].head(10)
print(sample_feedbacks.to_string(index=False))

save_simulation(df, output_filename)
print(f"\n Dataset saved to: {output_filename}")

print("\nBasic numeric stats:")

```

```
print(df[['monthly_credit','monthly_debit','ending_balance','monthly_invest_return_pct']].
      describe())
```

## 0.2 Part 2.1: Exploratory Data Analysis (EDA)

```
[2]: # EDA: load the saved CSV and perform quick exploratory analysis with visuals and summaries
pd.set_option('display.max_columns', None)

# Load dataset saved earlier (falls back to generated `df` if file missing)
try:
    df_eda = pd.read_csv(output_filename)
    print(f"Loaded dataset from {output_filename} - rows: {len(df_eda)}")
except Exception as e:
    print(f"Could not load {output_filename} ({e}), using in-memory `df` from simulation - rows: {len(df)}")
df_eda = df.copy()
```

Loaded dataset from data/sim\_customers\_seed\_42.csv - rows: 500

```
[15]: # Quick peek
print('\n-- Head (first 8 rows) --')
print(df_eda.head(8).to_string(index=False))

# Data types & missingness
print('\n-- Info & dtypes --')
print(df_eda.dtypes)

print('\n-- Missing values per column --')
print(df_eda.isnull().sum())

# Numeric summary
numeric_cols = [c for c in df_eda.columns if df_eda[c].dtype in
                ['int64','float64']]
print('\n-- Numeric summary (describe) --')
print(df_eda[numeric_cols].describe())

# Show value counts for categorical fields
cat_cols = ['occupation', 'transaction_category_major']
for c in cat_cols:
    if c in df_eda.columns:
        print(f"\nTop categories for {c}:")
        print(df_eda[c].value_counts(dropna=False).head(10))

# Visualizations (seaborn-only)
sns.set_theme(style='whitegrid')
```

```

# 1) Missingness map
ax = sns.heatmap(df_eda.isnull(), cbar=False)
ax.set_title('Missing values map (True = missing)')
ax.set_xlabel('Rows (index)')
ax.set_ylabel('Columns')
fig = ax.get_figure()
fig.set_size_inches(10,4)
fig.tight_layout()
plt.show()

# 2) Distributions (income, credit, debit, ending balance) - separate figures
sns.histplot(df_eda['monthly_income'].dropna(), bins=40, kde=True).
    ↪set(title='Monthly income')
plt.show()

sns.histplot(df_eda['monthly_credit'].dropna(), bins=40, kde=True).
    ↪set(title='Monthly credit')
plt.show()

sns.histplot(np.log1p(df_eda['monthly_debit'].dropna()), bins=40, kde=True).
    ↪set(title='Log(1 + monthly_debit)')
plt.show()

sns.histplot(df_eda['ending_balance'].dropna(), bins=40, kde=True).
    ↪set(title='Ending balance')
plt.show()

# 3) Boxplots by occupation for monthly_debit (clip extreme outliers for readability)
if 'occupation' in df_eda.columns:
    tmp = df_eda.copy()
    tmp['monthly_debit_clipped'] = tmp['monthly_debit'].
    ↪clip(upper=tmp['monthly_debit'].quantile(0.95))
    ax = sns.boxplot(x='occupation', y='monthly_debit_clipped', data=tmp)
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
    ax.set_title('Monthly debit by occupation (95th-percentile clipped)')
    fig = ax.get_figure()
    fig.tight_layout()
    plt.show()

# 4) Category counts
if 'transaction_category_major' in df_eda.columns:
    vc = df_eda['transaction_category_major'].value_counts(dropna=False)
    ax = sns.barplot(x=vc.values, y=vc.index)
    ax.set_title('Major transaction category (counts)')
    ax.set_xlabel('Count')

```

```

    ax.set_ylabel('Category')
    fig = ax.get_figure()
    fig.tight_layout()
    plt.show()

# 5) Correlation heatmap of numeric fields (auto-filtered for clean plot)
# Remove columns with zero variance or >50% missing values
num_for_corr = [
    c for c in numeric_cols
    if df_eda[c].nunique(dropna=True) > 1 and df_eda[c].notnull().sum() > len(df_eda)*0.5
]
if num_for_corr:
    corr = df_eda[num_for_corr].corr()
    ax = sns.heatmap(corr, annot=True, fmt=' .2f', cmap='coolwarm', center=0)
    ax.set_title('Correlation matrix (numeric features, cleaned)')
    fig = ax.get_figure()
    fig.set_size_inches(10,8)
    fig.tight_layout()
    plt.show()
else:
    print('No suitable numeric columns for correlation matrix after filtering.')

# 6) Small textual analysis: sample customer feedbacks
print('\n-- Sample (non-null) customer feedbacks (10) --')
if 'customer_feedback' in df_eda.columns:
    nonnull_fb = df_eda[df_eda['customer_feedback'].
    ↪notnull()][['customer_id', 'customer_feedback']]
    if lennonnull_fb == 0:
        print('No non-null feedbacks to show.')
    else:
        sample_fb = nonnull_fb.sample(n=min(10, lennonnull_fb),
        ↪random_state=42)
        print(sample_fb.to_string(index=False))
else:
    print('No customer_feedback column found.')

# Save a small EDA summary file for quick reference
eda_summary = {
    'rows': len(df_eda),
    'columns': list(df_eda.columns),
    'missing_per_column': df_eda.isnull().sum().to_dict(),
    'numeric_summary': df_eda[numeric_cols].describe().to_dict()
}

with open('data/eda_summary.json', 'w', encoding='utf-8') as f:
    json.dump(eda_summary, f, indent=2)

```

```
print('\n EDA complete - summary written to data/eda_summary.json')
```

Loaded dataset from data/sim\_customers\_seed\_42.csv - rows: 500

-- Head (first 8 rows) --

customer_id	date	age	occupation	monthly_income	monthly_credit	
monthly_debit	starting_balance	ending_balance	monthly_invest_gain	monthly_invest_return_pct	initial_credit_score	final_credit_score
transaction_category_major	customer_feedback					

CUST00001	2025-01-31	45	Teacher	3723.47	3801.07
2467.52	11200.47	12534.02		77.60	
0.52	668	673			dining

Dining was enjoyable, but finances remain strong.

CUST00002	2025-01-31	21	Teacher	5275.91	6245.50
6829.58	12190.17	11606.10		315.29	
1.74	633	638			healthcare

Healthcare costs were high, good reserves helped.

CUST00003	2025-01-31	45	Self-employed	3481.94	4089.60
5530.05	3511.02	2070.57		59.57	
0.79	696	686			dining

Dining out was excessive; need to cut back.

CUST00004	2025-01-31	50	Self-employed	3940.55	5116.70
6419.16	7142.85	5840.39		265.08	
2.21	748	753			utilities

Utilities were costly; exploring ways to save energy.

CUST00005	2025-01-31	32	Retail	5704.87	6097.01
2152.20	10235.67	14180.49		295.09	
1.84	705	700			dining

Excellent savings this month, very happy with progress!

CUST00006	2025-01-31	26	Engineer	6563.84	7271.12
5903.91	7811.20	9178.41		352.05	
2.39	627	622			dining

Managed dining costs well, financial health is good.

CUST00007	2025-01-31	36	Healthcare	3990.87	6456.65
8676.66	7461.86	5241.84		236.70	
1.73	743	738			shopping

Shopping spree impacted balance; will be more mindful.

CUST00008	2025-01-31	44	Finance	7823.06	8723.14
7201.27	24270.85	25792.72		494.00	
1.52	759	769			education

Education investment continues; strong savings too.

-- Info & dtypes --

customer_id	object
date	object
age	int64

```

occupation                      object
monthly_income                  float64
monthly_credit                  float64
monthly_debit                  float64
starting_balance                float64
ending_balance                 float64
monthly_invest_gain             float64
monthly_invest_return_pct       float64
initial_credit_score            int64
final_credit_score              int64
transaction_category_major     object
customer_feedback               object
dtype: object

-- Missing values per column --
customer_id                     0
date                            0
age                             0
occupation                      0
monthly_income                  10
monthly_credit                  15
monthly_debit                  15
starting_balance                0
ending_balance                 0
monthly_invest_gain             0
monthly_invest_return_pct       0
initial_credit_score            0
final_credit_score              0
transaction_category_major     0
customer_feedback               25
dtype: int64

-- Numeric summary (describe) --
      age   monthly_income  monthly_credit  monthly_debit \
count  500.00000      490.000000      485.000000      485.000000
mean   39.12200      3866.868020     4757.023814     6235.750087
std    11.06671      1977.735479     2564.605983     10977.072664
min    18.00000      500.000000      510.090000      717.820000
25%   32.00000      2401.985000     2917.520000     2456.830000
50%   39.00000      3813.220000     4598.850000     4189.850000
75%   47.00000      5304.980000     6245.500000     6629.330000
max   72.00000      9768.460000     15613.120000    148981.884457

      starting_balance  ending_balance  monthly_invest_gain \
count      500.00000      500.000000      500.000000
mean     7899.50754      7640.055400     127.123260
std     6291.82864      7246.865008     162.981478
min    -66.84000     -6347.860000      0.000000

```

25%	2977.16000	1855.445000	0.000000
50%	6270.91500	6041.695000	69.545000
75%	11205.29250	11323.280000	186.392500
max	31466.28000	35952.210000	1495.780000

	monthly_invest_return_pct	initial_credit_score	final_credit_score
count	500.000000	500.000000	500.000000
mean	0.857640	679.618000	678.118000
std	1.256535	48.601773	66.190743
min	-3.960000	506.000000	177.000000
25%	-0.002500	645.000000	644.000000
50%	0.845000	679.500000	681.000000
75%	1.690000	714.000000	714.000000
max	4.680000	837.000000	1038.000000

Top categories for occupation:

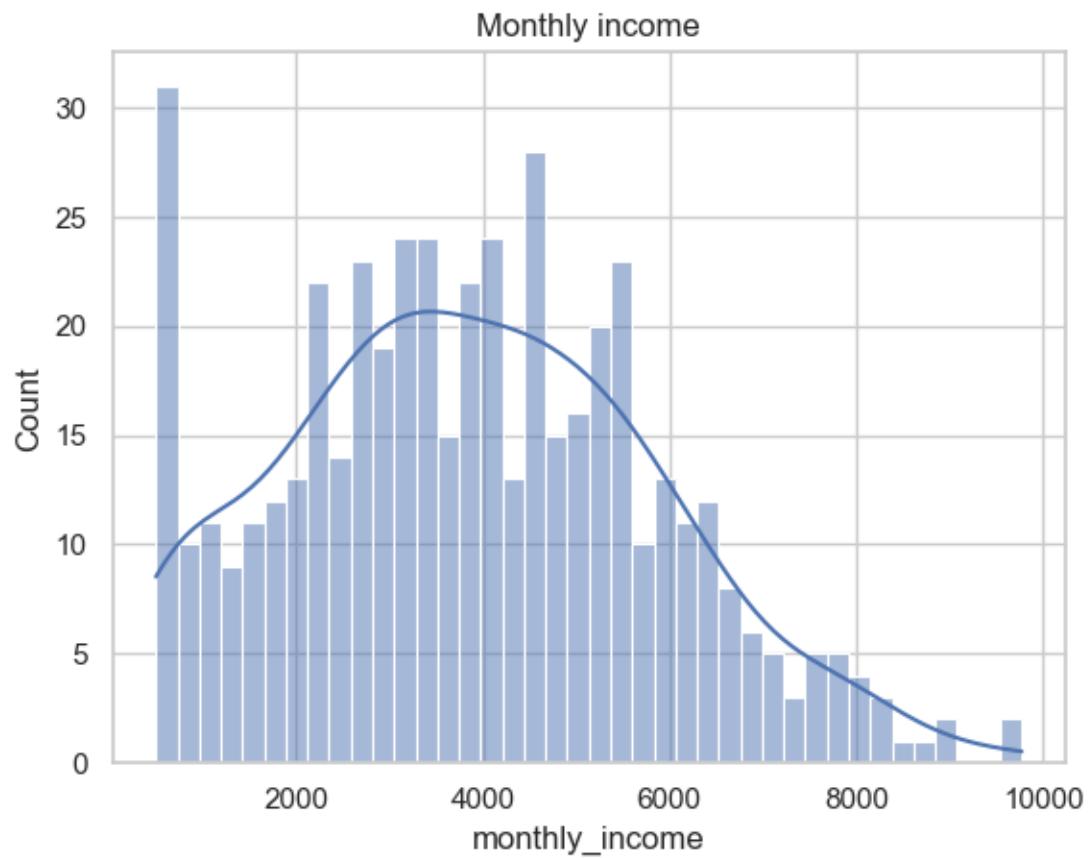
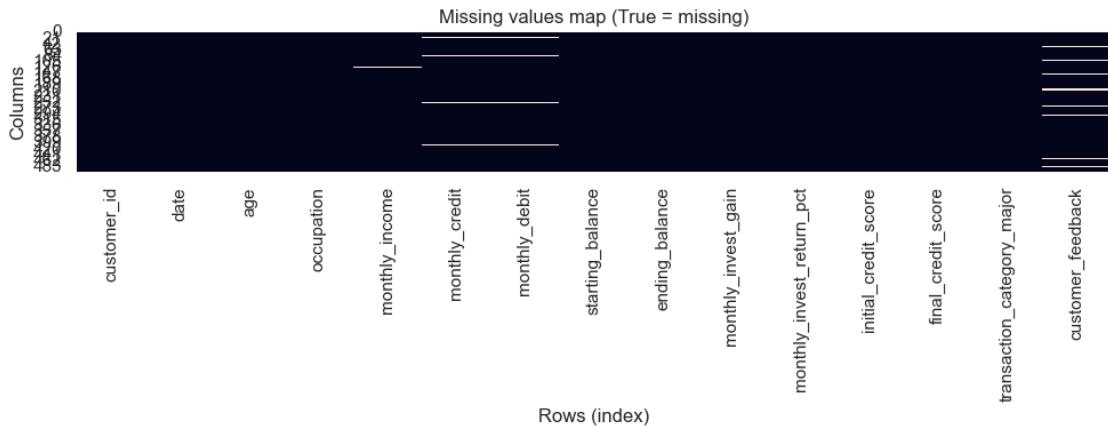
occupation	count
Self-employed	62
Retail	61
Manager	58
Teacher	54
Finance	49
Engineer	46
Student	44
Healthcare	43
Unemployed	43
Technician	40

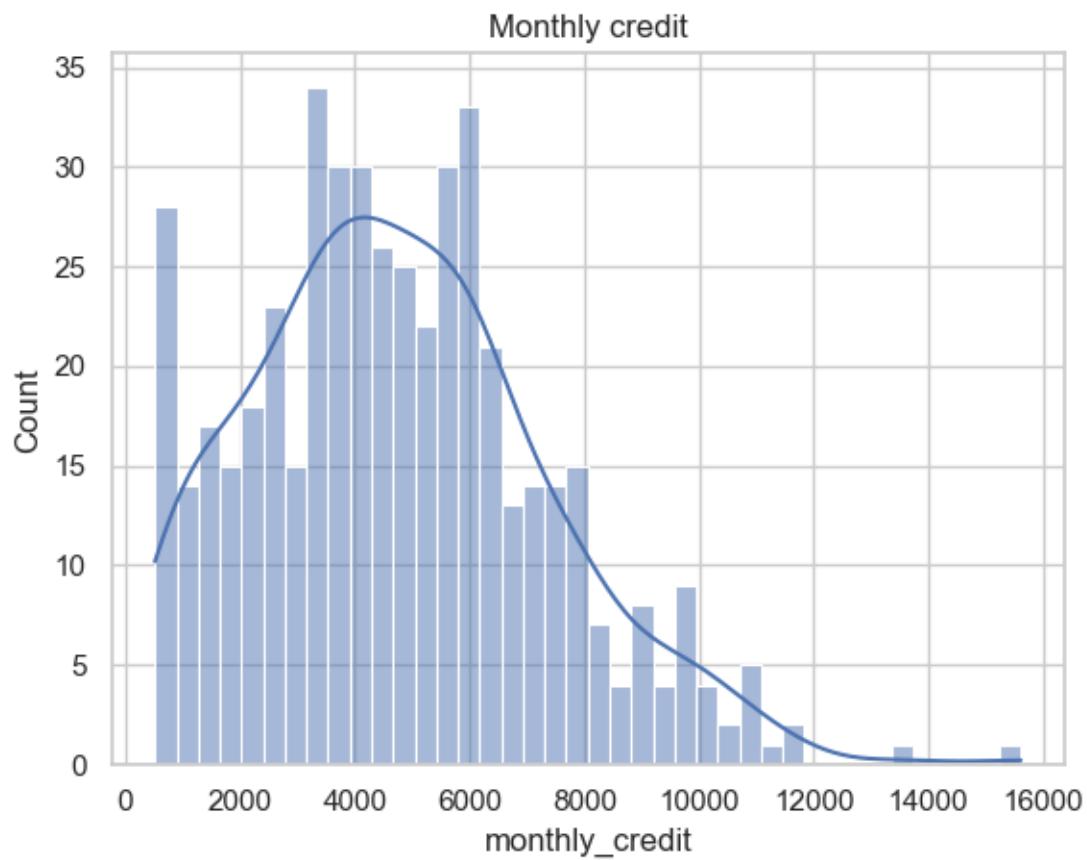
Name: count, dtype: int64

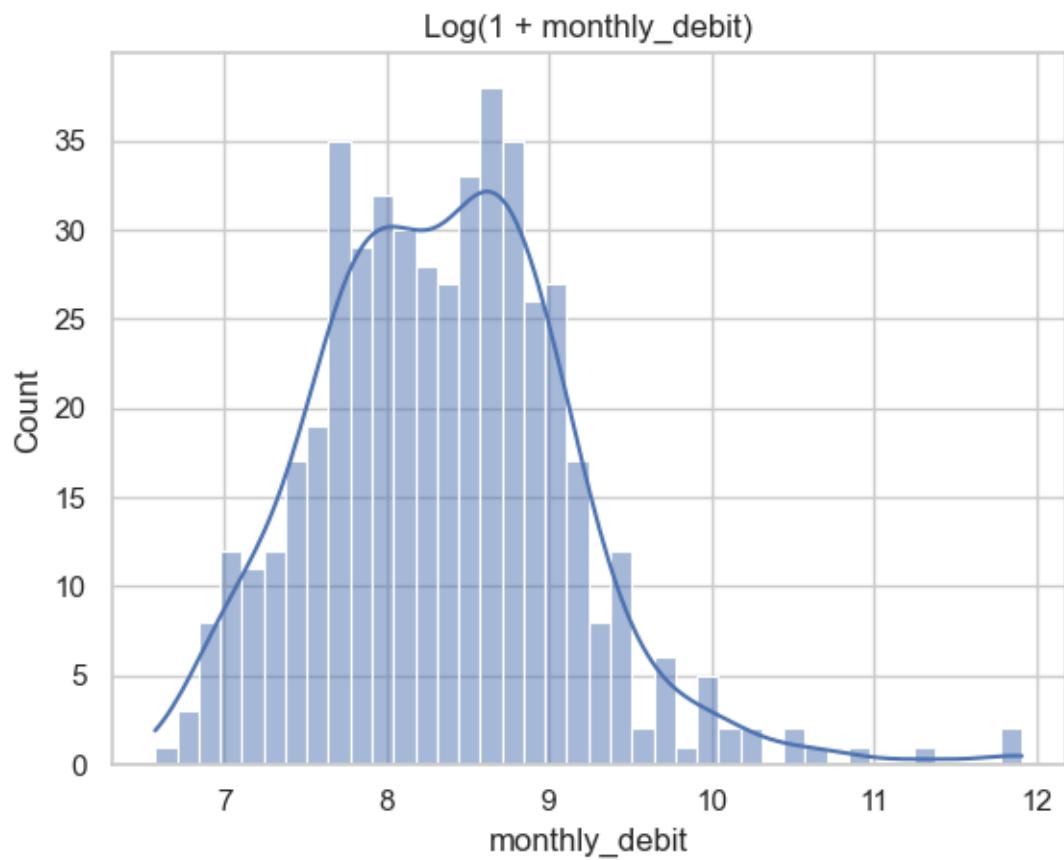
Top categories for transaction\_category\_major:

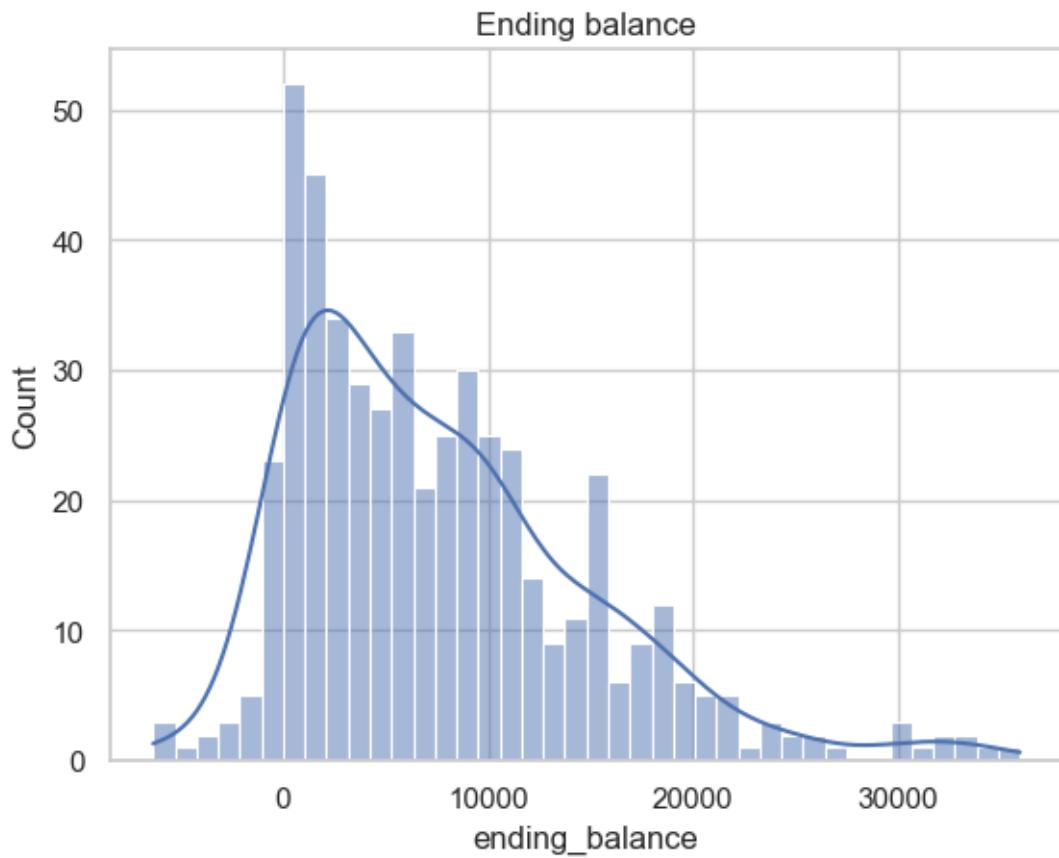
transaction_category_major	count
shopping	100
education	86
healthcare	83
entertainment	51
utilities	47
dining	46
groceries	43
transport	41
transfer	3

Name: count, dtype: int64

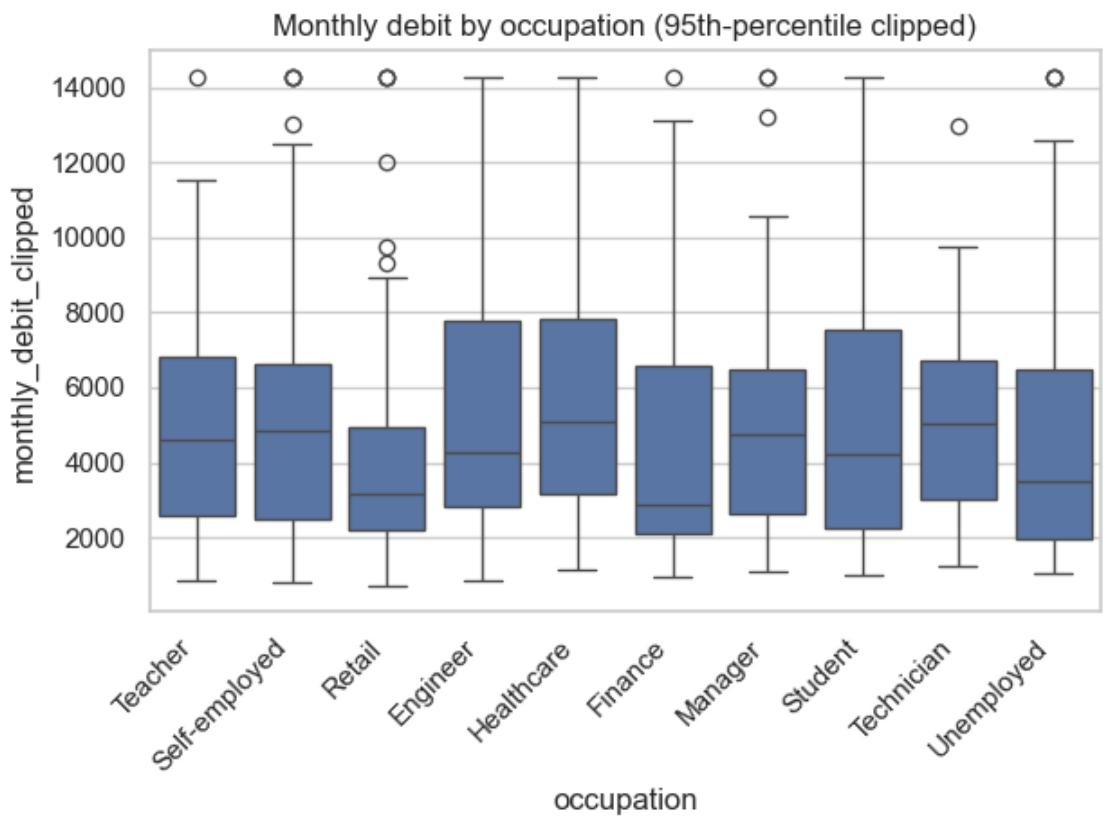


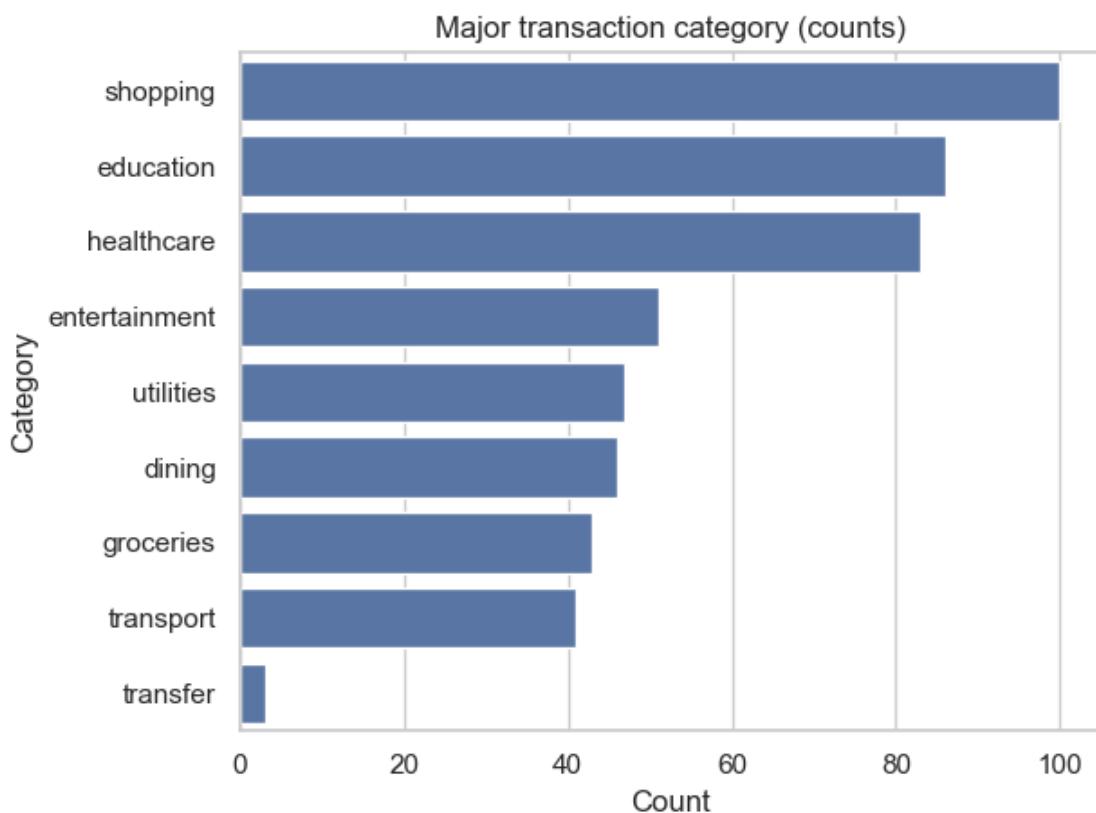


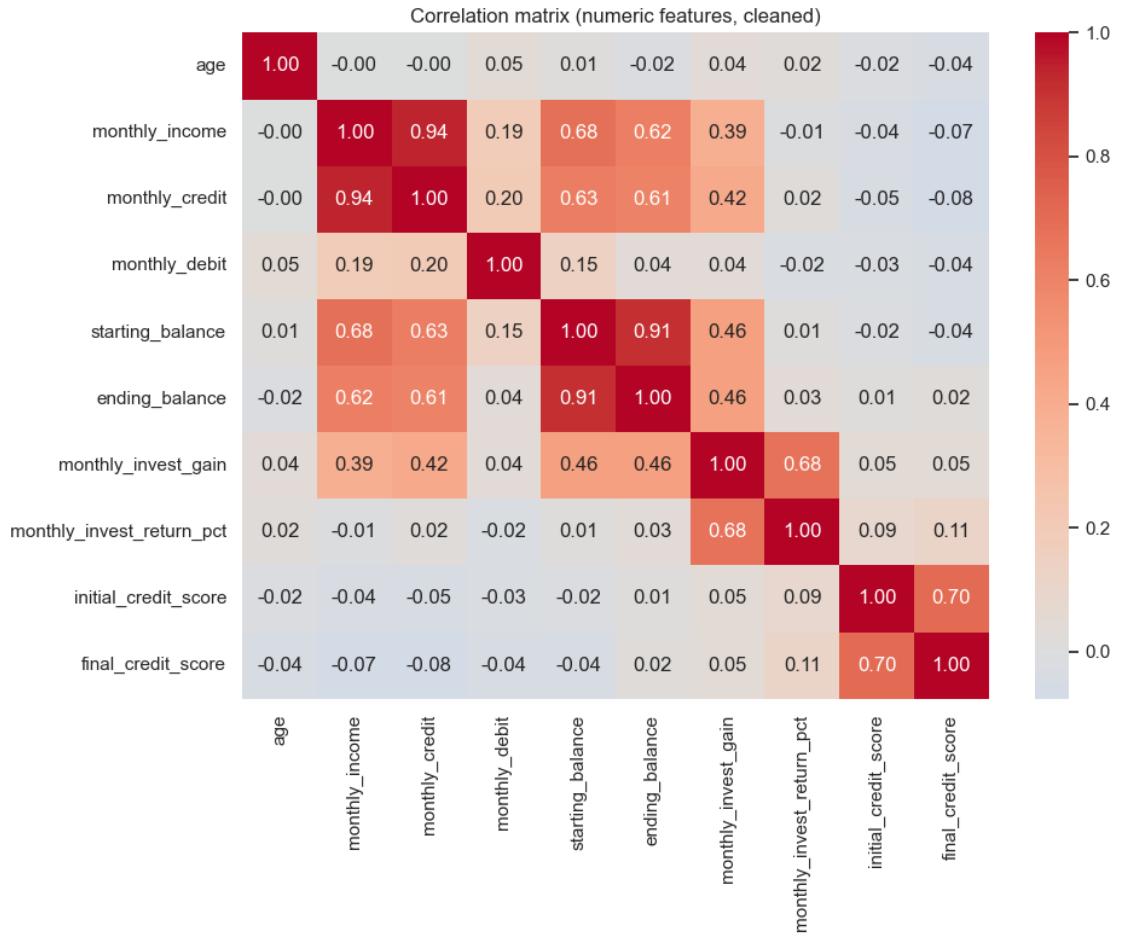




```
C:\Users\Jonas\AppData\Local\Temp\ipykernel_10444\2943935619.py:66: UserWarning:  
set_ticklabels() should only be used with a fixed number of ticks, i.e. after  
set_ticks() or using a FixedLocator.  
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
```







```
-- Sample (non-null) customer feedbacks (10) --
customer_id                                     customer_feedback
CUST00396           Debits slightly up. Keep an eye on healthcare spending.
CUST00444           Great savings! Keep up excellent management of grocery expenses.
CUST00010           Groceries were pricey; next month requires strict budgeting.
CUST00078           Excellent cash flow. Consider ways to optimize your utility bills.
CUST00375           High healthcare debits, but strong end balance. Good.
CUST00034           Shopping was okay; steady progress on financial goals.
CUST00083           Monitor dining expenses closely to prevent future balance drops.
CUST00465           Excellent month! Consider adding more to your long-term savings.
CUST00124           Good balance. Keep monitoring healthcare expenses carefully.
CUST00095           Monitor shopping closely to maintain your current healthy balance.
```

EDA complete - summary written to data/eda\_summary.json

[11]: # Additional charts

```

# Age distribution
sns.histplot(df_eda['age'].dropna(), bins=20, kde=True)
plt.title('Customer Age Distribution')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()

# Credit score change distribution
if 'final_credit_score' in df_eda.columns and 'initial_credit_score' in df_eda.
    ↵columns:
    score_change = df_eda['final_credit_score'] - df_eda['initial_credit_score']
    sns.histplot(score_change.dropna(), bins=20, kde=True)
    plt.title('Credit Score Change Distribution')
    plt.xlabel('Final - Initial Credit Score')
    plt.ylabel('Count')
    plt.show()

# Scatter: monthly_income vs. monthly_debit
sns.scatterplot(x='monthly_income', y='monthly_debit', data=df_eda)
plt.title('Monthly Income vs. Monthly Debit')
plt.xlabel('Monthly Income')
plt.ylabel('Monthly Debit')
plt.show()

# Boxplot: ending_balance by occupation
if 'occupation' in df_eda.columns:
    plt.figure(figsize=(10,5))
    sns.boxplot(x='occupation', y='ending_balance', data=df_eda)
    plt.title('Ending Balance by Occupation')
    plt.xlabel('Occupation')
    plt.ylabel('Ending Balance')
    plt.xticks(rotation=45, ha='right')
    plt.tight_layout()
    plt.show()

# Scatter: monthly_invest_return_pct vs. final_credit_score
if 'monthly_invest_return_pct' in df_eda.columns and 'final_credit_score' in df_eda.
    ↵columns:
    sns.scatterplot(x='monthly_invest_return_pct', y='final_credit_score', u
        ↵data=df_eda)
    plt.title('Investment Return % vs. Final Credit Score')
    plt.xlabel('Monthly Investment Return (%)')
    plt.ylabel('Final Credit Score')
    plt.show()

# Outlier highlighting for ending_balance and monthly_debit (top/bottom 1%)
def highlight_outliers(series, quantile=0.01):

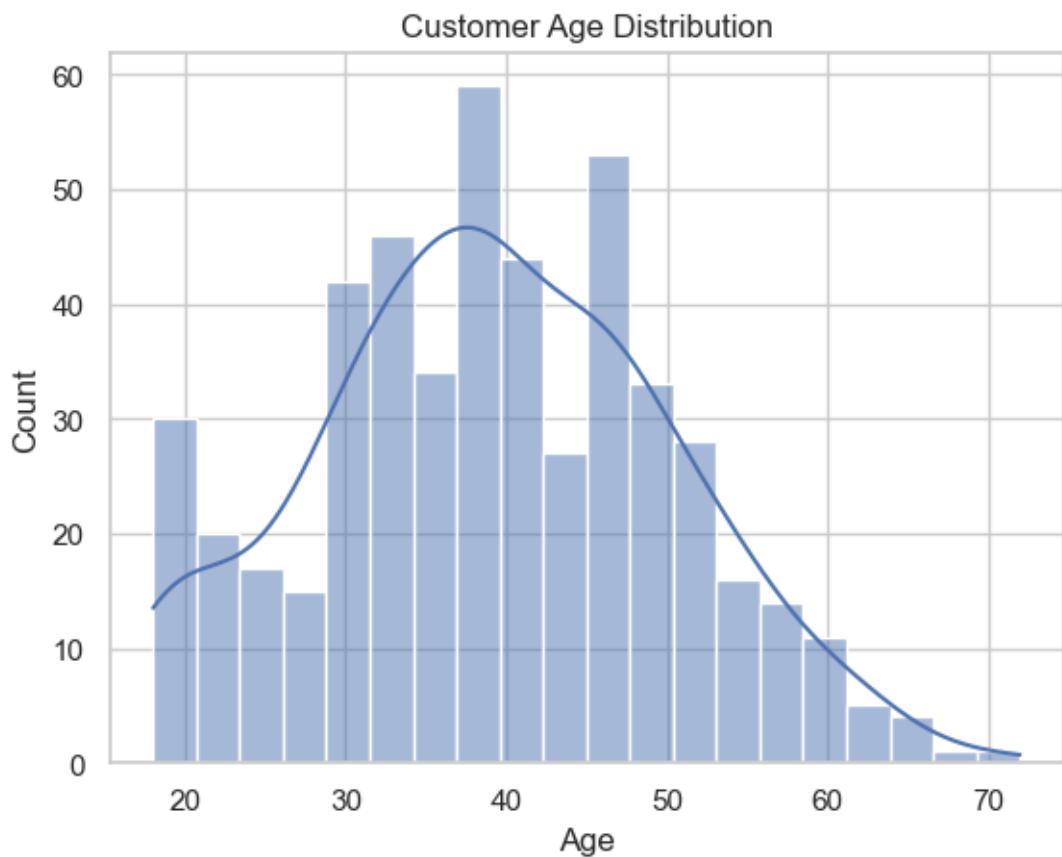
```

```

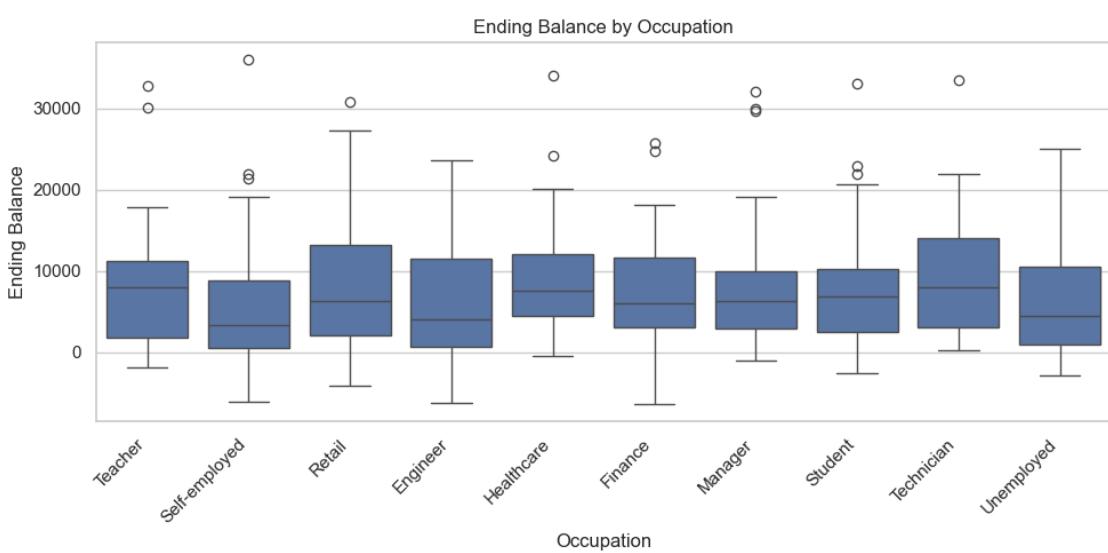
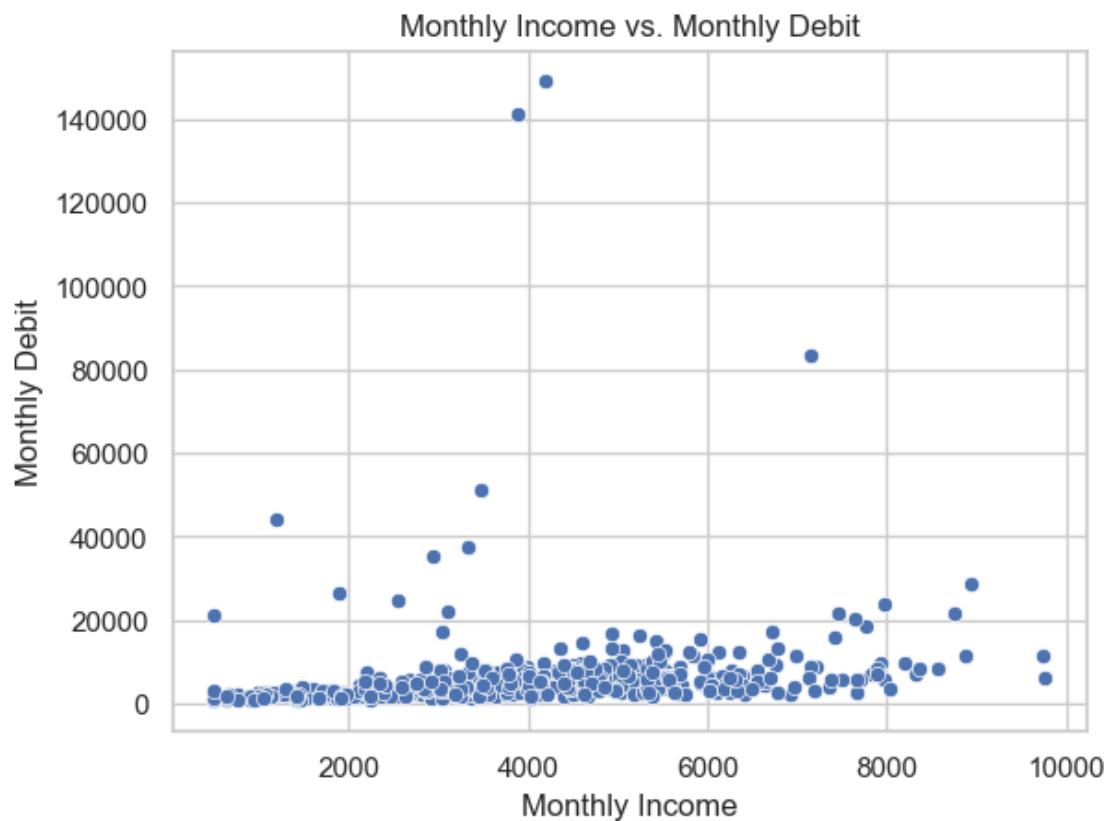
lower = series.quantile(quantile)
upper = series.quantile(1-quantile)
return (series < lower) | (series > upper)

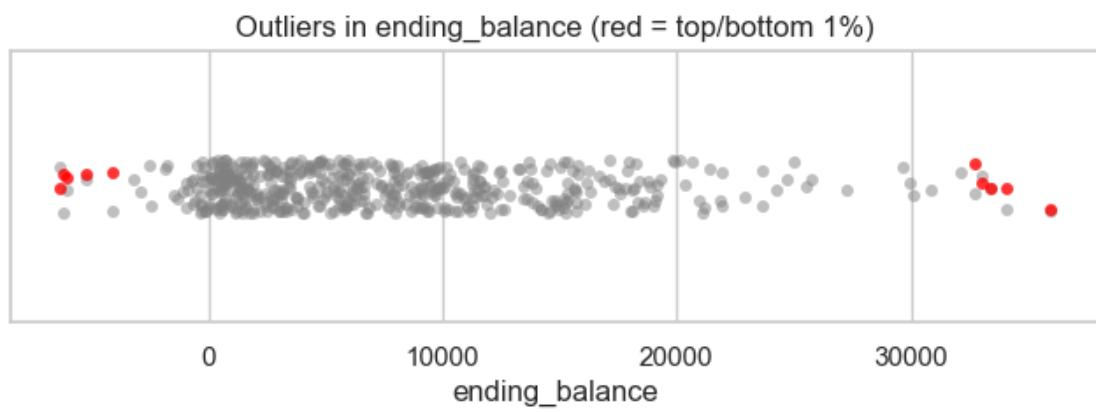
for col in ['ending_balance', 'monthly_debit']:
    if col in df_eda.columns:
        outlier_mask = highlight_outliers(df_eda[col])
        plt.figure(figsize=(8,2))
        sns.stripplot(x=df_eda[col], color='gray', alpha=0.5)
        sns.stripplot(x=df_eda.loc[outlier_mask, col], color='red', alpha=0.8)
        plt.title(f'Outliers in {col} (red = top/bottom 1%)')
        plt.xlabel(col)
        plt.show()

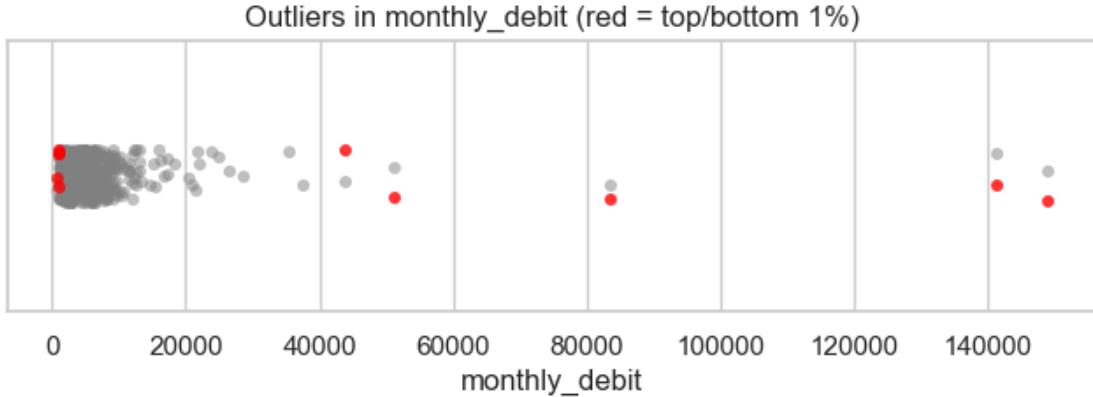
```











### 0.3 Part 2.2: AI-generated summaries

```
[3]: # This cell provides consolidated EDA helpers and an AI-assisted summarizer.
```

```
from pathlib import Path

OUT_DIR = Path('data')
OUT_DIR.mkdir(parents=True, exist_ok=True)

def summary_statistics(df: pd.DataFrame):
    numeric_cols = df.select_dtypes(include=["number"]).columns.tolist()
    cat_cols = df.select_dtypes(include=["object"]).columns.tolist()
    stats = {"rows": len(df), "numeric": {}, "categorical_top": {}}
    if numeric_cols:
        stats['numeric'] = df[numeric_cols].describe().to_dict()
    for c in cat_cols:
        stats['categorical_top'][c] = df[c].value_counts(dropna=False).head(10).to_dict()
    return stats

def detect_anomalies(df: pd.DataFrame, numeric_cols: list) -> pd.DataFrame:
    if not numeric_cols:
        return pd.DataFrame()
    num = df[numeric_cols].copy()
    mu = num.mean()
    sigma = num.std(ddof=0).replace(0, np.nan)
    z = (num - mu) / sigma
    z_abs = z.abs()
    z_mask = (z_abs > 3).any(axis=1)

    iqr_mask = pd.Series(False, index=df.index)
```

```

for c in numeric_cols:
    col = df[c]
    if col.dropna().empty:
        continue
    q1 = col.quantile(0.25); q3 = col.quantile(0.75)
    iqr = q3 - q1
    low = q1 - 1.5 * iqr; high = q3 + 1.5 * iqr
    iqr_mask = iqr_mask | ((col < low) | (col > high))

combined = z_mask | iqr_mask
anomalies = df[combined].copy()
if not anomalies.empty:
    anomalies['_max_abs_z'] = z_abs.loc[anomalies.index].max(axis=1)
return anomalies

def generate_ai_summaries(stats, corr_matrix, anomalies, df):
    # Build a short prompt-like summary for Gemini to consume
    lines = [f"rows: {stats.get('rows')}"]
    numeric = stats.get('numeric', {})
    for col in [
        'monthly_income', 'monthly_credit', 'monthly_debit', 'ending_balance', 'monthly_invest_return',
    ]:
        try:
            m = numeric[col]['mean']
            lines.append(f"mean_{col}: {m:.2f}")
        except Exception:
            pass

    corr_lines = []
    try:
        if corr_matrix is not None:
            pairs = []
            cols = corr_matrix.columns.tolist()
            for i,a in enumerate(cols):
                for j,b in enumerate(cols):
                    if i>=j: continue
                    v = corr_matrix.iloc[i,j]
                    if abs(v) >= 0.4:
                        pairs.append((a,b,float(v)))
            pairs = sorted(pairs, key=lambda x: -abs(x[2]))
            for a,b,v in pairs[:6]:
                corr_lines.append(f"{a} vs {b}: {v:.2f}")
    except Exception:
        pass

```

```

prompt_body = "\n".join(lines + ["--- correlations ---"] + corr_lines +_
↪[f"anomalies_count: {len(anomalies)}"])

# Primary path: use Google GenAI (Gemini) if SDK and API key are present
if 'genai' in globals() and genai is not None and 'types' in globals() and_
↪types is not None and os.getenv('GOOGLE_API_KEY'):

    try:
        client = genai.Client(api_key=os.getenv('GOOGLE_API_KEY'))
        system = (
            "You are a data analyst. Based on the stats and correlations_"
↪provided, produce up to 6 short plain-English insights."
            "Return a JSON array of objects with keys 'topic' and 'insight'."
↪ Each insight should be 1-2 sentences."
        )
        contents = [system + "\n\n" + prompt_body]
        response = client.models.generate_content(
            model="gemini-2.5-flash",
            contents=contents,
            config=types.GenerateContentConfig(system_instruction=system)
        )
        text = _extract_text(response)
        # Extract JSON substring if the model wraps the output
        start = text.find('[')
        end = text.rfind(']')
        if start != -1 and end != -1 and end > start:
            json_text = text[start:end+1]
            parsed = json.loads(json_text)
            out_lines = []
            for o in parsed[:6]:
                topic = o.get('topic') or o.get('title') or 'insight'
                insight = o.get('insight') or o.get('text') or o.
↪get('feedback') or ''
                out_lines.append(f"- {topic}: {insight}")
            return "\n".join(out_lines)
        # fallback: return raw text if JSON not found
        return text
    except Exception as e:
        print(f"GenAI call failed: {e} - falling back to heuristic_"
↪summarizer")

# Fallback heuristic (used only if GenAI not available)
insights = []
try:
    if 'monthly_income' in df.columns and 'monthly_invest_return_pct' in df.
↪columns:

```

```

        c = df[['monthly_income','monthly_invest_return_pct']].dropna().corr().iloc[0,1]
        if math.isnan(c): c = 0.0
        if c>0.2:
            insights.append('Higher income weakly associates with higher investment return pct.')
        elif c<-0.2:
            insights.append('Higher income tends to have lower investment return pct in this sample.')
        else:
            insights.append('No strong relationship between income and investment return pct.')
    except Exception:
        pass

    try:
        if 'monthly_income' in df.columns and 'ending_balance' in df.columns:
            c = df[['monthly_income','ending_balance']].dropna().corr().iloc[0,1]
            if c>0.5:
                insights.append('Monthly income strongly correlates with ending balance.')
            elif c>0.2:
                insights.append('Monthly income moderately correlates with ending balance.')
            else:
                insights.append('Weak relationship between income and ending balance.')
    except Exception:
        pass

    try:
        if 'monthly_debit' in df.columns and 'ending_balance' in df.columns:
            c = df[['monthly_debit','ending_balance']].dropna().corr().iloc[0,1]
            if c<-0.5:
                insights.append('Higher debits strongly associate with lower ending balances.')
            elif c<-0.2:
                insights.append('Debits moderately negatively associate with ending balance.')
            else:
                insights.append('No strong negative relationship between debits and ending balance.')
    except Exception:
        pass

```

```

try:
    if 'transaction_category_major' in df.columns:
        top = df['transaction_category_major'].value_counts().head(3).
        ↪to_dict()
        top_str = ', '.join([f'{k}({v})' for k,v in top.items()])
        insights.append(f"Top categories: {top_str}.")
except Exception:
    pass

try:
    if len(anomalies)==0:
        insights.append('Very few extreme outliers detected across numeric
        ↪fields.')
    else:
        s = anomalies.head(3)
        insights.append(f"Detected {len(anomalies)} anomalous rows. Sample
        ↪IDs: {', '.join(map(str, s.get('customer_id', s.index).tolist()[:3]))} .")
except Exception:
    pass

try:
    if 'initial_credit_score' in df.columns and 'final_credit_score' in df.
    ↪columns:
        changes = (df['final_credit_score'] - df['initial_credit_score']).
        ↪dropna()
        if not changes.empty:
            mean_change = changes.mean()
            if mean_change>2:
                insights.append('Credit scores increased slightly on
        ↪average.')
            elif mean_change<-2:
                insights.append('Credit scores decreased slightly on
        ↪average.')
            else:
                insights.append('Average credit score change is small.')
except Exception:
    pass

insights = insights[:6]
return '\n'.join([f"- {s}" for s in insights])

```

```

[7]: # Run the consolidated EDA + AI summary on the dataframe `df_eda` if present
      ↪(falls back to `df`)
try:
    working_df = df_eda
except NameError:

```

```

working_df = globals().get('df', None)

if working_df is None:
    raise RuntimeError('Could not find a dataframe to run EDA on (no df_eda or
        ↵df in notebook scope.)')

stats = summary_statistics(working_df)
with open(OUT_DIR / 'eda_combined_stats.json', 'w', encoding='utf-8') as f:
    json.dump(stats, f, indent=2, default=lambda x: str(x))

numeric_cols = working_df.select_dtypes(include=['number']).columns.tolist()
# correlation
corr = None
if len(numeric_cols) >= 2:
    corr = working_df[numeric_cols].corr()
    # save a CSV of correlations for quick inspection
    corr.to_csv(OUT_DIR / 'correlations.csv')

# anomalies
anoms = detect_anomalies(working_df, numeric_cols)
anoms.to_csv(OUT_DIR / 'anomalies_combined.csv', index=False)

# AI or heuristic summary
ai_text = generate_ai_summaries(stats, corr, anoms, working_df)
with open(OUT_DIR / 'ai_summary.txt', 'w', encoding='utf-8') as f:
    f.write(ai_text)

print('\n==== Data Anomalies ====\n')
if anoms.empty:
    print('No anomalies detected.')
else:
    print(anoms.head(10).to_string(index=False))

print('\n==== AI / heuristic insights ====\n')
print(ai_text)
print(f"\nWrote combined EDA outputs to: {OUT_DIR.resolve()}")

```

==== Data Anomalies ====

customer_id	date	age	occupation	monthly_income	monthly_credit
monthly_debit	starting_balance	ending_balance	monthly_invest_gain		
monthly_invest_return_pct	initial_credit_score	final_credit_score			
transaction_category_major					
customer_feedback	_max_abs_z				
CUST00008	2025-01-31	44	Finance	7823.06	8723.14
7201.270000		24270.85		25792.72	494.00
1.52		759		769	education

Education investment continues; strong savings too. 2.604607

CUST00011	2025-01-31	55	Student	3110.72	4170.34
22035.915826		2123.53		2634.51	121.59
1.97		728		718	dining
Dining was a treat, but aiming for higher savings.					1.440865
CUST00012	2025-01-31	22	Manager	6930.78	8292.94
2132.040000		23506.19		29667.09	0.00
-0.40		731		741	groceries
Groceries managed well; excellent progress on savings!					3.042570
CUST00021	2025-01-31	36	Student	6495.48	8196.25
3628.920000		17371.65		21938.99	833.55
3.37		638		633	education
Education expenses met, excellent financial position overall.					4.338740
CUST00023	2025-01-31	45	Finance	8577.36	9518.93
8448.620000		15233.93		16304.24	482.97
1.99		673		668	education
Education expenses are significant, but manageable with income.					2.384194
CUST00036	2025-01-31	55	Finance	4359.44	7750.06
13107.640000		10704.38		5346.80	58.87
0.32		656		646	entertainment
Entertainment costs very high, significantly impacted balance.					1.436190
CUST00040	2025-01-31	23	Unemployed	4748.96	7150.13
2980.340000		9902.44		14072.24	614.77
3.74		699		689	education
Education expenses met; strong savings continue this month.					2.995034
CUST00045	2025-01-31	27	Manager	6963.33	8055.31
4041.690000		25977.81		29991.42	270.10
0.80		778		783	groceries
Groceries managed; phenomenal savings, very well done!					3.087369
CUST00047	2025-01-31	44	Finance	4936.63	5946.68
16739.800000		4445.26		-6347.86	0.00
-0.01		645		640	healthcare
Serious overspending, especially on healthcare. Immediate action needed.					
1.932135					
CUST00049	2025-01-31	35	Self-employed	7775.38	9735.37
18380.080000		8361.95		-282.76	0.00
-1.06		704		694	entertainment
Entertainment spending excessive, resulted in an overdrawn account.					1.978276

== AI / heuristic insights ==

- Income and Credit Activity: There is a very strong positive correlation (0.94) between monthly income and monthly credit activity. This indicates that higher-income individuals tend to have substantially more credit flowing into their accounts.
- Account Balance Stability: An individual's starting balance is a powerful predictor of their ending balance, with a correlation of 0.91. This suggests that account balances are relatively stable or follow a consistent trend over

the period observed.

- Data Quality/Anomalies: A significant 15.2% of the records contain anomalies (76 out of 500 rows). This high number warrants further investigation into potential data quality issues or unusual financial behaviors.
- Spending vs. Income: Average monthly debits (6235.75) significantly exceed both average monthly income (3866.87) and credit (4757.02). This suggests that many individuals are spending more than their current income and credit inflows, potentially relying on savings or other financial resources.
- Credit Score Predictability: An individual's initial credit score is a strong predictor of their final credit score, with a correlation of 0.70. While some change is possible, credit scores tend to show considerable stability over time.
- Income's Impact on Balances: Higher monthly income is strongly correlated with a higher starting account balance (0.68). This indicates that a greater income generally contributes to a stronger foundational financial position.

Wrote combined EDA outputs to: C:\Users\Jonas\OneDrive - Universiti Malaya\UM2025-26\code\wie3007\_indiv\data

#### 0.4 Part 3: Data-preprocessing

```
[32]: # Use Gemini for sentiment and topic classification
def generate_sentiment_topic_gemini_batch(feedbacks, n_requests=10,
                                         model="gemini-2.5-flash-lite", max_retries=3, retry_delay=5,
                                         batch_delay_seconds=4):
    """Send feedbacks to Gemini in batches and return sentiment and topic maps.

    feedbacks: list of dicts, each with 'customer_id' and 'feedback'.
    Returns: {customer_id: {'sentiment': str, 'topic': str}}
    """
    if 'genai' not in globals() or genai is None or 'types' not in globals() or
       types is None:
        raise RuntimeError("Google GenAI SDK not available.")

    GOOGLE_API_KEY = os.getenv("GOOGLE_API_KEY")
    if not GOOGLE_API_KEY:
        raise ValueError("GOOGLE_API_KEY not found.")

    client = genai.Client(api_key=GOOGLE_API_KEY)

    result_map = {}
    total = len(feedbacks)
    if total == 0:
        return result_map

    batch_size = max(1, math.ceil(total / n_requests))

    for i in range(0, total, batch_size):
```

```

batch = feedbacks[i:i+batch_size]
items_text = []
for f in batch:
    items_text.append(f"ID: {f['customer_id']} | Feedback:{f['feedback']}")
prompt_body = "\n".join(items_text)
prompt = (
    "You are given multiple customer feedbacks (one per line in the\n"
    "format 'ID: <id> | Feedback: <text>').\n"
    "For each feedback, classify the sentiment as 'positive',\n"
    "'neutral', or 'negative', and detect the topic as one of: 'savings', 'loan',\n"
    "'investment', 'credit', 'dining', 'healthcare', 'utilities', 'shopping',\n"
    "'education', 'transport', 'groceries', 'entertainment'.\n"
    "Return a JSON array of objects with fields {\"customer_id\": <id>,\n"
    "\"sentiment\": <sentiment>, \"topic\": <topic>} and NOTHING ELSE.\n"
    "Respond with valid JSON only.\n\n"
    f"Feedbacks:\n{prompt_body}"
)

attempt = 0
resp_text = None
last_exception = None
while attempt < max_retries:
    try:
        print(f"Sending Gemini batch request for feedbacks {i} to\n"
              f"{i+len(batch)-1}...")
        response = client.models.generate_content(
            model=model,
            contents=[prompt],
            config=types.
        GenerateContentConfig(system_instruction="Classify sentiment and topic for\n"
                                         "customer feedbacks.")
    )
        resp_text = _extract_text(response)
        if resp_text:
            break
        else:
            raise RuntimeError('Empty response from Gemini')
    except Exception as e:
        last_exception = e
        attempt += 1
        wait = retry_delay * (2 ** (attempt - 1))
        print(f"Gemini batch request failed (attempt {attempt}/\n"
              f"{max_retries}): {e}. Retrying in {wait}s...")
        time.sleep(wait)

```

```

    if resp_text is None:
        raise RuntimeError(f"Gemini failed for batch starting at index {i} after {max_retries} attempts.") from last_exception

    if resp_text.startswith('```json') and resp_text.endswith('```'):
        resp_text = resp_text[7:-3].strip()
    elif resp_text.startswith('```') and resp_text.endswith('```'):
        resp_text = resp_text[3:-3].strip()
    try:
        parsed = json.loads(resp_text)
    except Exception:
        start = resp_text.find('[')
        end = resp_text.rfind(']')
        if start != -1 and end != -1:
            parsed = json.loads(resp_text[start:end+1])

    if not isinstance(parsed, list):
        raise ValueError(f"Invalid response for batch {i}.")

    for obj in parsed:
        cid = obj.get('customer_id')
        sent = obj.get('sentiment')
        top = obj.get('topic')
        if cid and sent and top:
            result_map[str(cid)] = {'sentiment': str(sent), 'topic': str(top)}

    time.sleep(batch_delay_seconds)

    return result_map

```

```

[33]: # Data Preprocessing using LLMs / SLMs

# Copy df_eda to df_preproc and drop ID columns
df_preproc = df_eda.copy()
df_preproc = df_preproc.drop(columns=['customer_id'])

# 1. Handle missing values
num_cols = df_preproc.select_dtypes(include=['number']).columns
cat_cols = df_preproc.select_dtypes(include=['object']).columns.tolist()

# Exclude text fields from categorical encoding
text_fields = ['customer_feedback']
cat_encode_cols = [col for col in cat_cols if col not in text_fields]

# Impute numerics with median
if len(num_cols) > 0:

```

```

    num_imputer = SimpleImputer(strategy='median')
    df_prepoc[num_cols] = num_imputer.fit_transform(df_prepoc[num_cols])

# Impute categoricals with mode
if len(cat_encode_cols) > 0:
    cat_imputer = SimpleImputer(strategy='most_frequent')
    df_prepoc[cat_encode_cols] = cat_imputer.
    ↪fit_transform(df_prepoc[cat_encode_cols])

# 2. Normalize numerical features
if len(num_cols) > 0:
    scaler = StandardScaler()
    df_prepoc[num_cols] = scaler.fit_transform(df_prepoc[num_cols])

# 3. Encode categorical variables (excluding text fields)
if len(cat_encode_cols) > 0:
    for col in cat_encode_cols:
        le = LabelEncoder()
        df_prepoc[col] = le.fit_transform(df_prepoc[col].astype(str))

# 4. Process text fields (customer_feedback)
# Impute missing feedback with 'No feedback provided'
if 'customer_feedback' in df_prepoc.columns:
    df_prepoc['customer_feedback'] = df_prepoc['customer_feedback'].
    ↪fillna('No feedback provided')

# Prepare feedbacks for Gemini
feedback_list = []
for idx, row in df_prepoc.iterrows():
    feedback_list.append({'customer_id': row.get('customer_id', idx), ↪
    'feedback': row['customer_feedback']})

# Get sentiments and topics from Gemini
st_map = generate_sentiment_topic_gemini_batch(feedback_list, ↪
n_requests=10, batch_delay_seconds=4)
df_prepoc['feedback_sentiment'] = df_prepoc.apply(lambda row: st_map.
get(str(row.get('customer_id', row.name)), {}).get('sentiment', 'neutral'), ↪
axis=1)
df_prepoc['feedback_topic'] = df_prepoc.apply(lambda row: st_map.
get(str(row.get('customer_id', row.name)), {}).get('topic', 'other'), axis=1)

print('Preprocessing complete. Sample:')
display(df_prepoc.head())

```

Sending Gemini batch request for feedbacks 0 to 49...

Sending Gemini batch request for feedbacks 50 to 99...

```

Sending Gemini batch request for feedbacks 100 to 149...
Sending Gemini batch request for feedbacks 150 to 199...
Sending Gemini batch request for feedbacks 200 to 249...
Sending Gemini batch request for feedbacks 250 to 299...
Sending Gemini batch request for feedbacks 300 to 349...
Sending Gemini batch request for feedbacks 350 to 399...
Sending Gemini batch request for feedbacks 400 to 449...
Sending Gemini batch request for feedbacks 450 to 499...
Preprocessing complete. Sample:

```

	date	age	occupation	monthly_income	monthly_credit	monthly_debit	\
0	0.0	0.531674	0.939857	-0.072768	-0.376958	-0.343048	
1	0.0	-1.639164	0.939857	0.720964	0.591754	0.060636	
2	0.0	0.531674	0.210154	-0.196258	-0.262615	-0.059628	
3	0.0	0.983932	0.210154	0.038221	0.144418	0.022654	
4	0.0	-0.644196	-0.154697	0.940283	0.532908	-0.372229	
	starting_balance	ending_balance	monthly_invest_gain	\			
0	0.525168	0.675998		-0.304163			
1	0.682625	0.547825		1.155685			
2	-0.698189	-0.769307		-0.414899			
3	-0.120381	-0.248586		0.847304			
4	0.371673	0.903423		1.031620			
	monthly_invest_return_pct	initial_credit_score	final_credit_score	\			
0	-0.268976	-0.239284		-0.077399			
1	0.702920	-0.960144		-0.606704			
2	-0.053884	0.337403		0.119199			
3	1.077339	1.408395		1.132439			
4	0.782584	0.522767		0.330921			
	transaction_category_major	customer_feedback	feedback_sentiment	\			
0	-1.511365	-1.152617		negative			
1	0.140400	-0.119693		neutral			
2	-1.511365	-1.173554		negative			
3	1.792165	1.471567		positive			
4	-1.511365	-0.719906		negative			
	feedback_topic						
0	dining						
1	dining						
2	dining						
3	dining						
4	dining						